Machine Learning Engineering Intern

Location: Kansas City, MO (Hybrid)

Compensation: Competitive (starting at $22)

About GenLift

GenLift is a fast-growing AI consulting firm specializing in deploying cutting-edge generative AI and retrieval-augmented systems for enterprise clients across healthcare, finance, and professional services. Our expert team delivers custom ML solutions that unlock insights from sensitive, proprietary data while ensuring compliance and security.

Role Overview

As an ML Engineering Intern at GenLift, you'll work closely with our students and researchers to design, implement, and evaluate machine learning workflows and pipelines. You'll gain hands-on experience with real client projects—ranging from fine-tuning large language models to building data ingestion and evaluation tools—and contribute directly to production-ready systems. You'll contribute to building a diverse set of applications and learn a lot through the process.

Key Responsibilities

- **Prototype Deployment:** Containerize and deploy model inference services using Docker and Kubernetes; integrate with RESTful APIs and orchestration tools.

- **Experimentation & Analysis:** Run hyperparameter sweeps; analyze results to recommend improvements; document findings in reproducible notebooks.

- **Collaboration & Documentation:** Work alongside cross-functional teams (data engineers, software developers, product managers) to translate business requirements into technical solutions; maintain clear, organized documentation.

Qualifications

- **Academic Standing:** Pursuing a Bachelor's or Master's degree in Computer Science, Data Science, Electrical Engineering, or a related field.

- **Technical Skills:**

    - Strong proficiency in Python and relevant ML libraries (PyTorch, scikit-learn).

    - Comfortable working with SQL and NoSQL databases.

- **Analytical Mindset:** Solid understanding of statistics, optimization, and evaluation metrics for classification, regression, and sequence generation tasks.

- **Communication:** Excellent written and verbal communication skills; able to present technical concepts to both technical and non-technical audiences.

- **Team Player:** Self-motivated, curious, and collaborative—eager to learn and contribute in a fast-paced environment.

Nice-to-Haves

- Prior internship or research experience in ML or data engineering.

- Familiarity with RAG (Retrieval-Augmented Generation) architectures and vector databases (e.g., Pinecone, FAISS), NLP frameworks (Hugging Face Transformers), Docker, Kubernetes, or similar container/orchestration tools.

- Hands-on experience with cloud platforms (AWS, Azure, or GCP).

- Exposure to CI/CD pipelines, MLOps tools, or MLflow.

What You'll Gain

- Direct mentorship from senior ML researchers and engineers.

- Exposure to end-to-end enterprise AI projects with high-impact outcomes.

- Opportunity to develop production-grade machine learning systems.

- A collaborative, supportive culture that values innovation and professional growth.

To Apply:

Please send your resume, a brief email highlighting relevant projects, and a link to any code samples or GitHub repositories to **jeet@genlift.io** with the subject line "ML Engineering Intern Application." We look forward to hearing how you can help us drive innovation at GenLift!